


Proton-MRI deployment pipeline



Deployed acceleration: shift-augmented sliding-window classifier · batched BEV WEPL engine · bucket-batched forward · streamed zlib output ($\sim 2\times$; native $1\times 1\times 3$ mm grids write tens of GB/job) · no torch.compile. Job runtime 414 s $\rightarrow$ 189 s (compression) $\rightarrow$ 110 s (batched engine).